

oueees-201506

Part 2: large- scale information systems

Kenji Rikitake

16-JUN-2015

School of Engineering Science

Osaka University

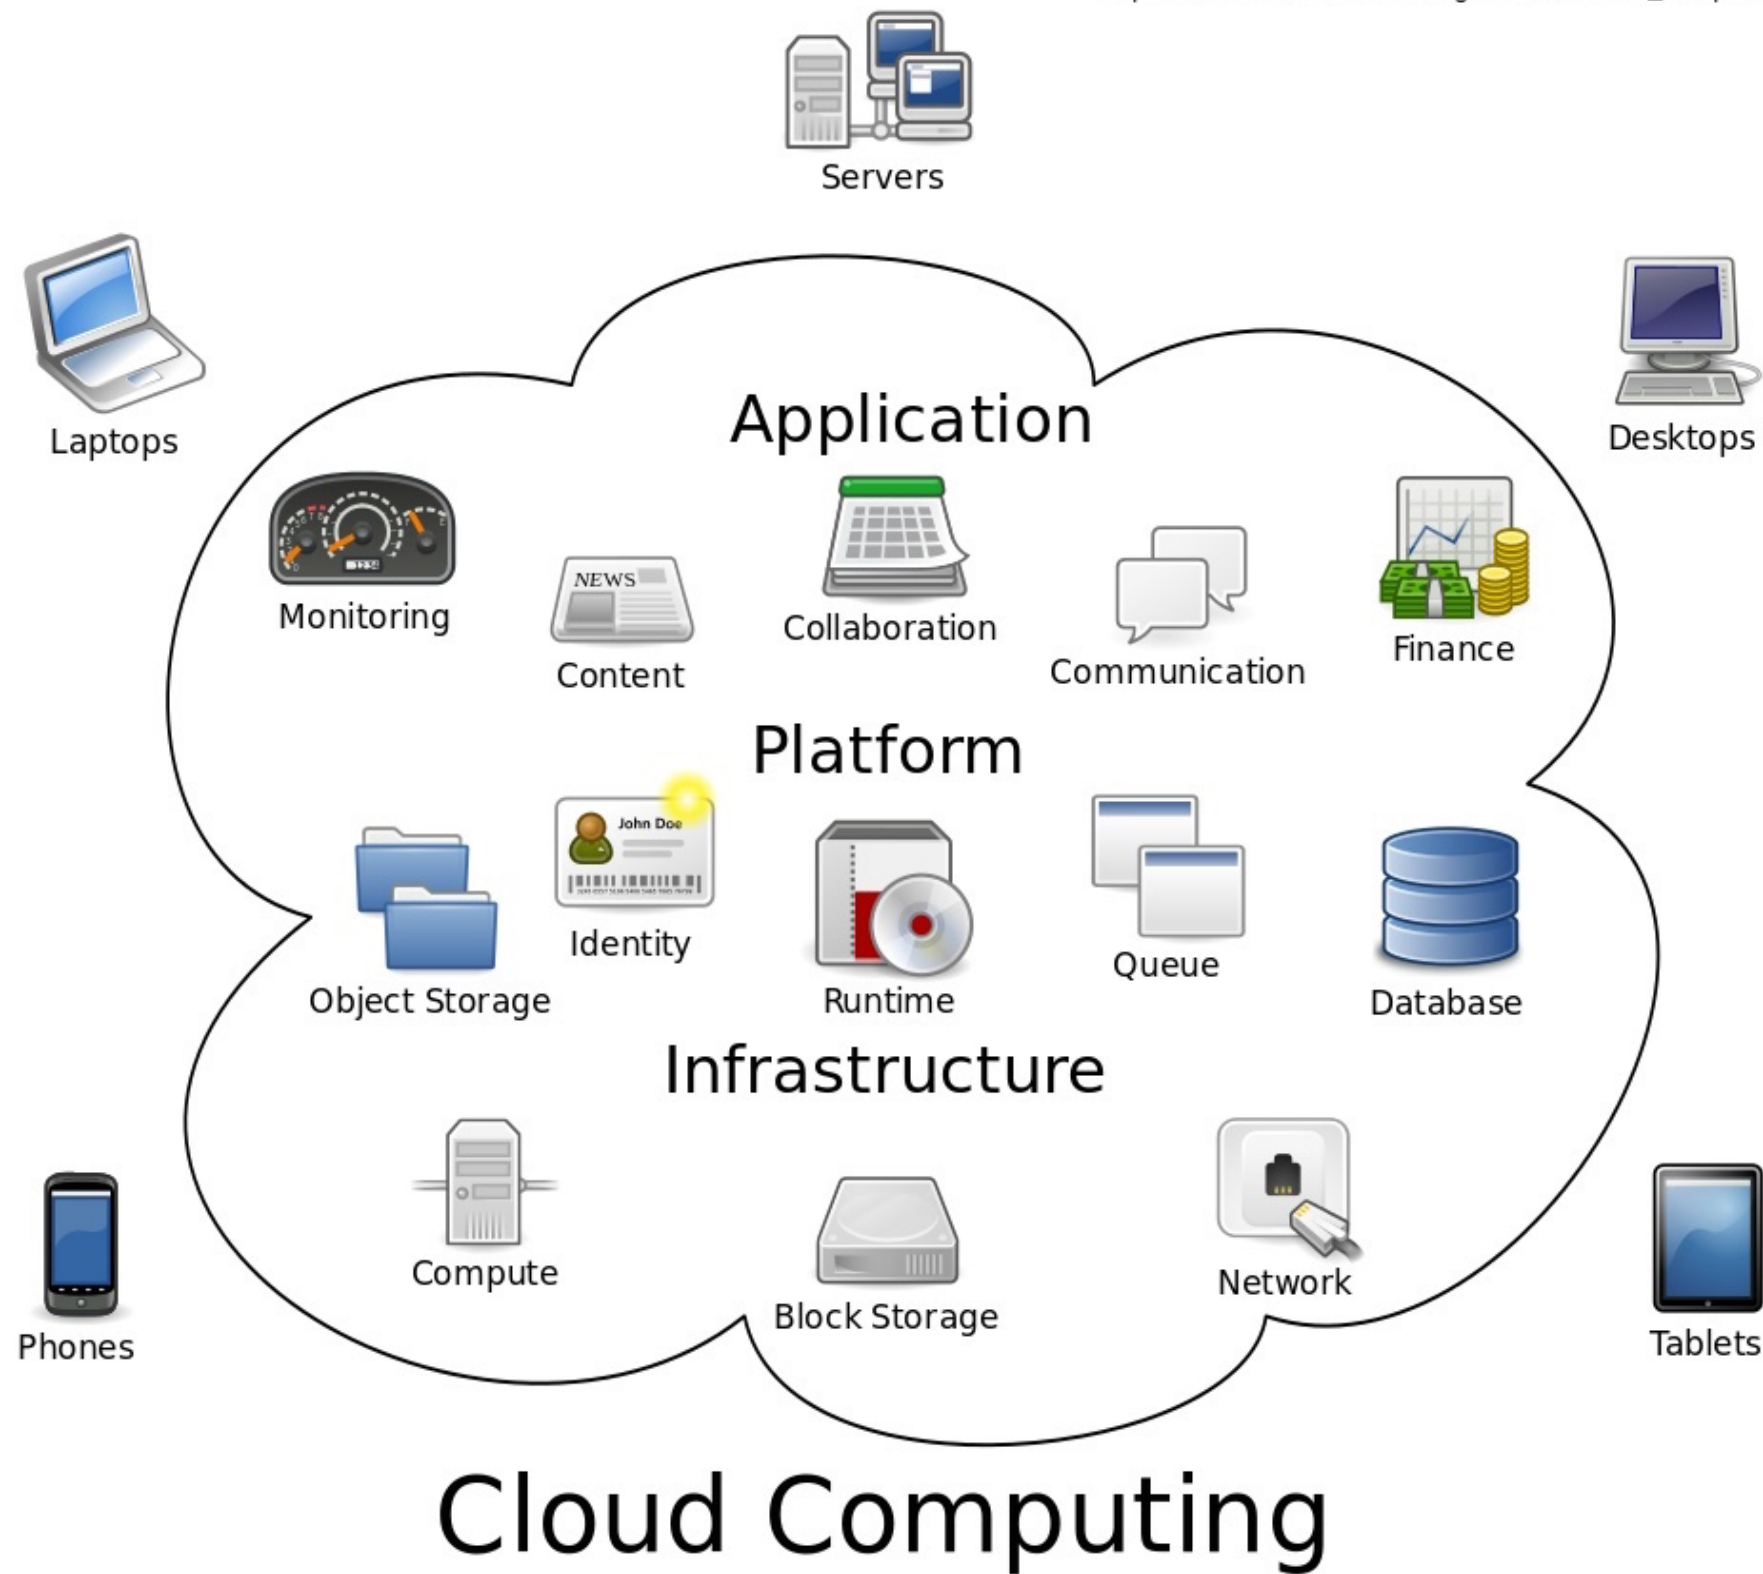
Toyonaka, Osaka, Japan

@jj1bdx

Lecture notes on GitHub

- <https://github.com/jj1bdx/oueees-201505-public/>
- Don't forget to *check out the issues!*

Cloud computing systems




Cloud computing elements

- Servers and services on the Internet
- Endpoint terminals (smartphones, tablets, laptops, etc.) *outside* the cloud
- *Highly centralized systems depending on the Internet*

Inside the services

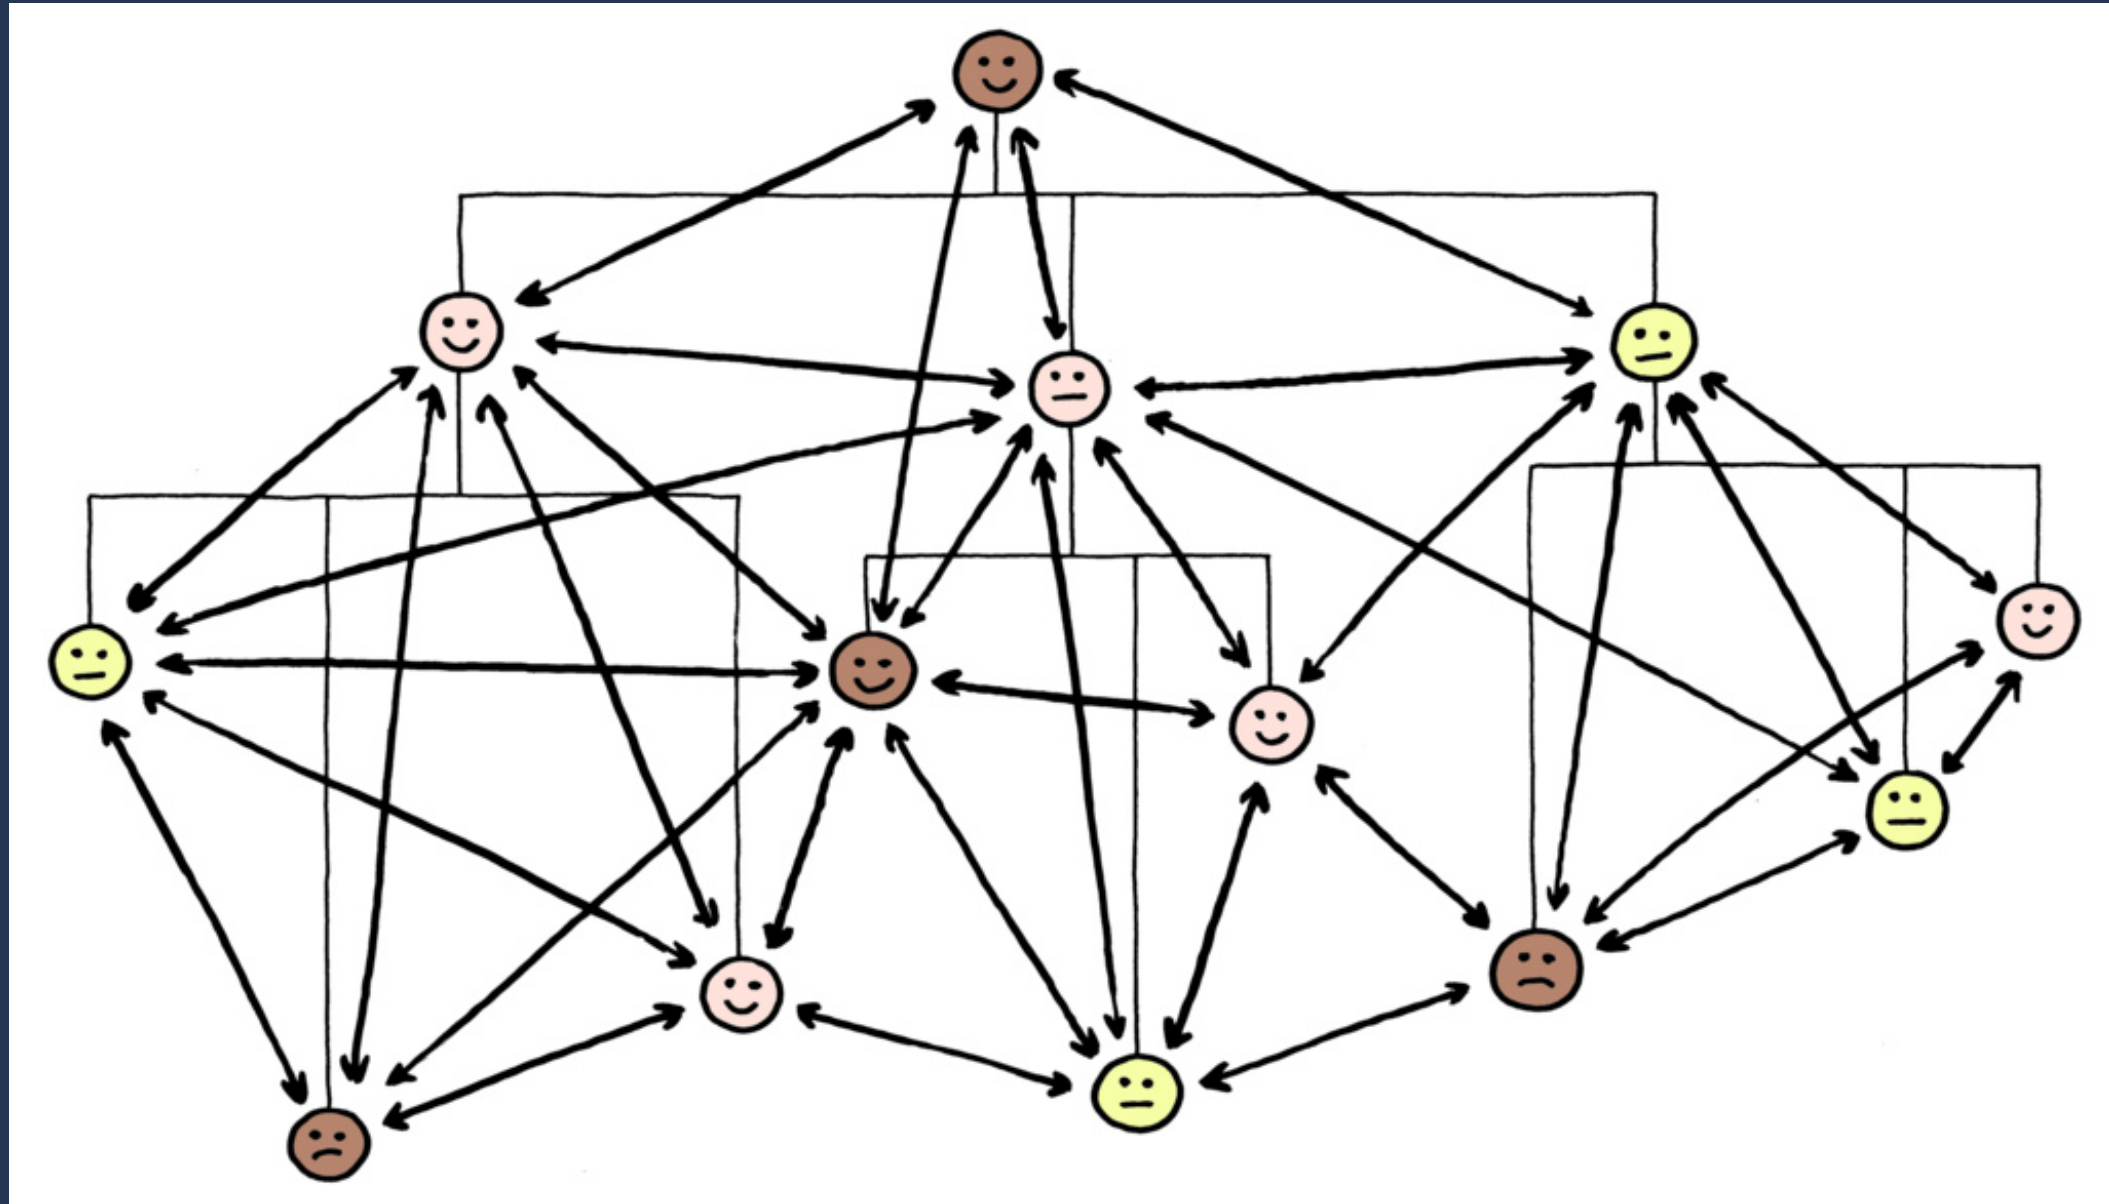
- A cluster of distributed systems
- Multiple computers *collaboratively* connected to do the same task
- *Highly decentralized or even distributed*

Forms of networks ¹



¹ Carl S. Sterner, Resilience and Decentralization, http://www.carlsterner.com/research/2009_resilience_and_decentralization.shtml

Real world: hierarchy *and* decentralization ²



² By Jurgen Appelo, licensed CC BY 2.0, <https://www.flickr.com/photos/jurgenappelo/5201869924/>

Centralized social behavior accelerated by cloud computing

- Sharing *everything* - no privacy
- *Panopticon*³ style of governance, filtering, censorship, or autocracy
- Complete *externalization* of resources, leading to *no personal control*

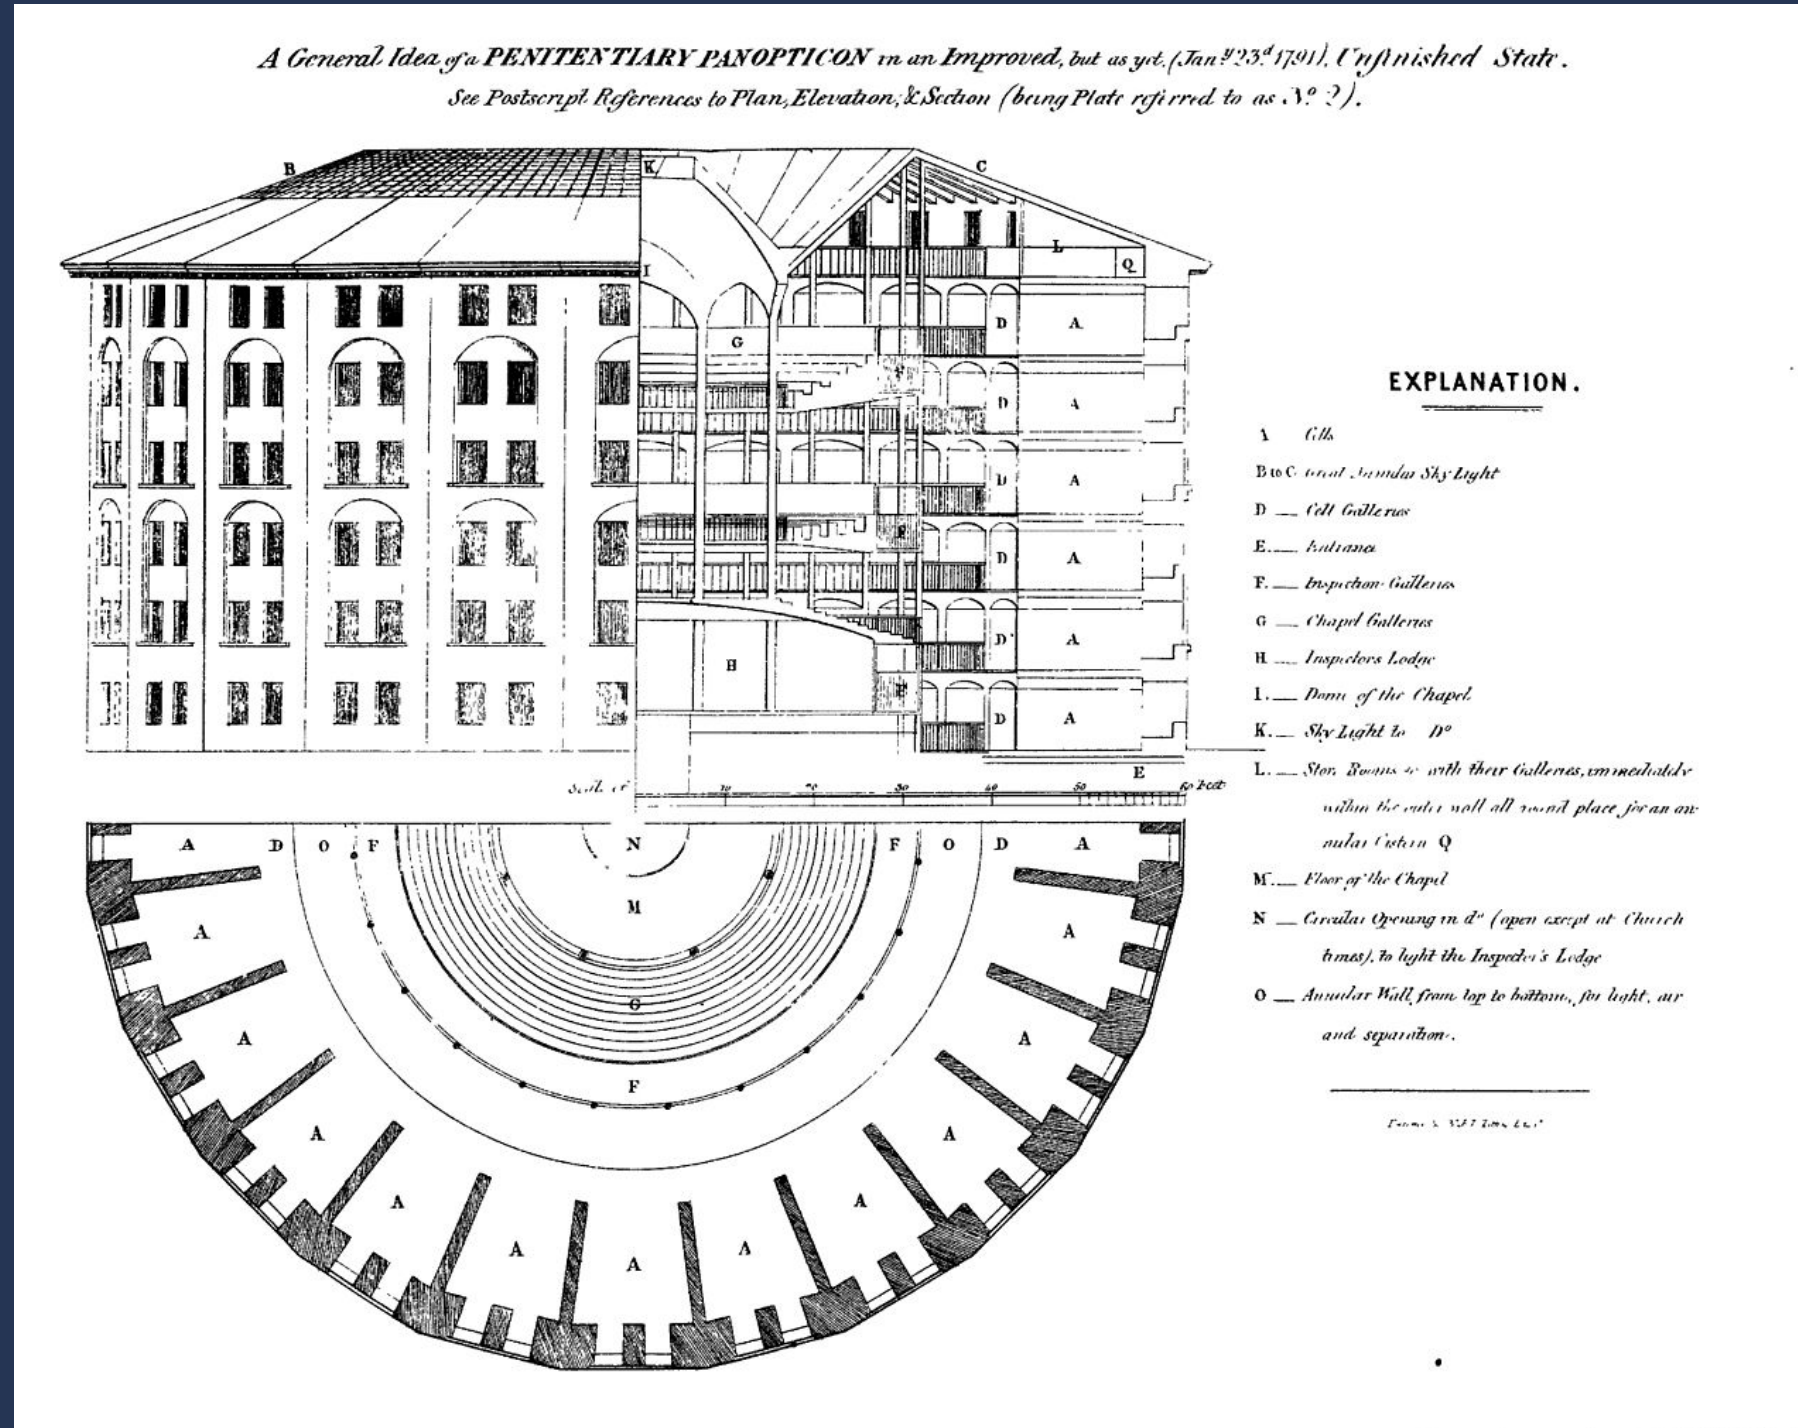
³ n. a circular prison with cells arranged around a central well, from which prisoners could at all times be observed. (New Oxford American Dictionary, Apple OS X 10.10.3)

Precidio Modelo Prison⁴



⁴ Friman, licensed CC BY-SA 3.0, <https://en.wikipedia.org/wiki/Panopticon#/media/File:Presidio-modelo2.JPG>

Panoption plan example (public domain)



INGSOC

The slogans: ⁵

- War is peace
- Freedom is slavery
- Ignorance is strength
- Independent thinking = *thoughtcrime*

NOTE: this is a *fiction*!

⁵ George Orwell, "Nineteen Eighty-Four", 1949.

Perpetual War 常在戦場

**Why cloud
computing has
become so
dystopian?**


**We have sold
freedom for
*convenience***

Convenience of centralized systems

- Ubiquitous/global accessibility
- Concentrated data for easy analysis
- Easy control of the information flow
- No extra cost for sharing
- No need to think about where the information locates

The inconvenient truth of centralized systems

What if the core/cloud fails?



Inconvenience of centralized systems

- Ubiquity or *no accessibility*
- When the core fails, no alternative
- When the core loses data, *no backup*
- The system performance is restricted by the capability of the core
- Endpoint systems will lost *all capabilities*

**Centralized
systems are *not*
sustainable**

Sustainable information systems: decentralized and distributed

Real-world challenges

- Natural disasters
- Device failures
- Human operation errors
- Political impediments
- Social resentments

Handling *failures*

- Redundancy: keeping backup units ready
- Fault tolerance: keeping systems running even the components fail
- Resilience by failing fast: early detection of failures and invocation of the recovery procedures


Why fault tolerance?

- Hard disk MTBF \approx 1 million hours
- 1000 hard disks running 24 hours x 365 days
= 8.76 million hours
- If you're running a system with 1000 hard disks, *nine* out of 1000 will fail in a year
- Recovery of a disk content takes often *a day*
- You can't stop a system for *a day*, can you?

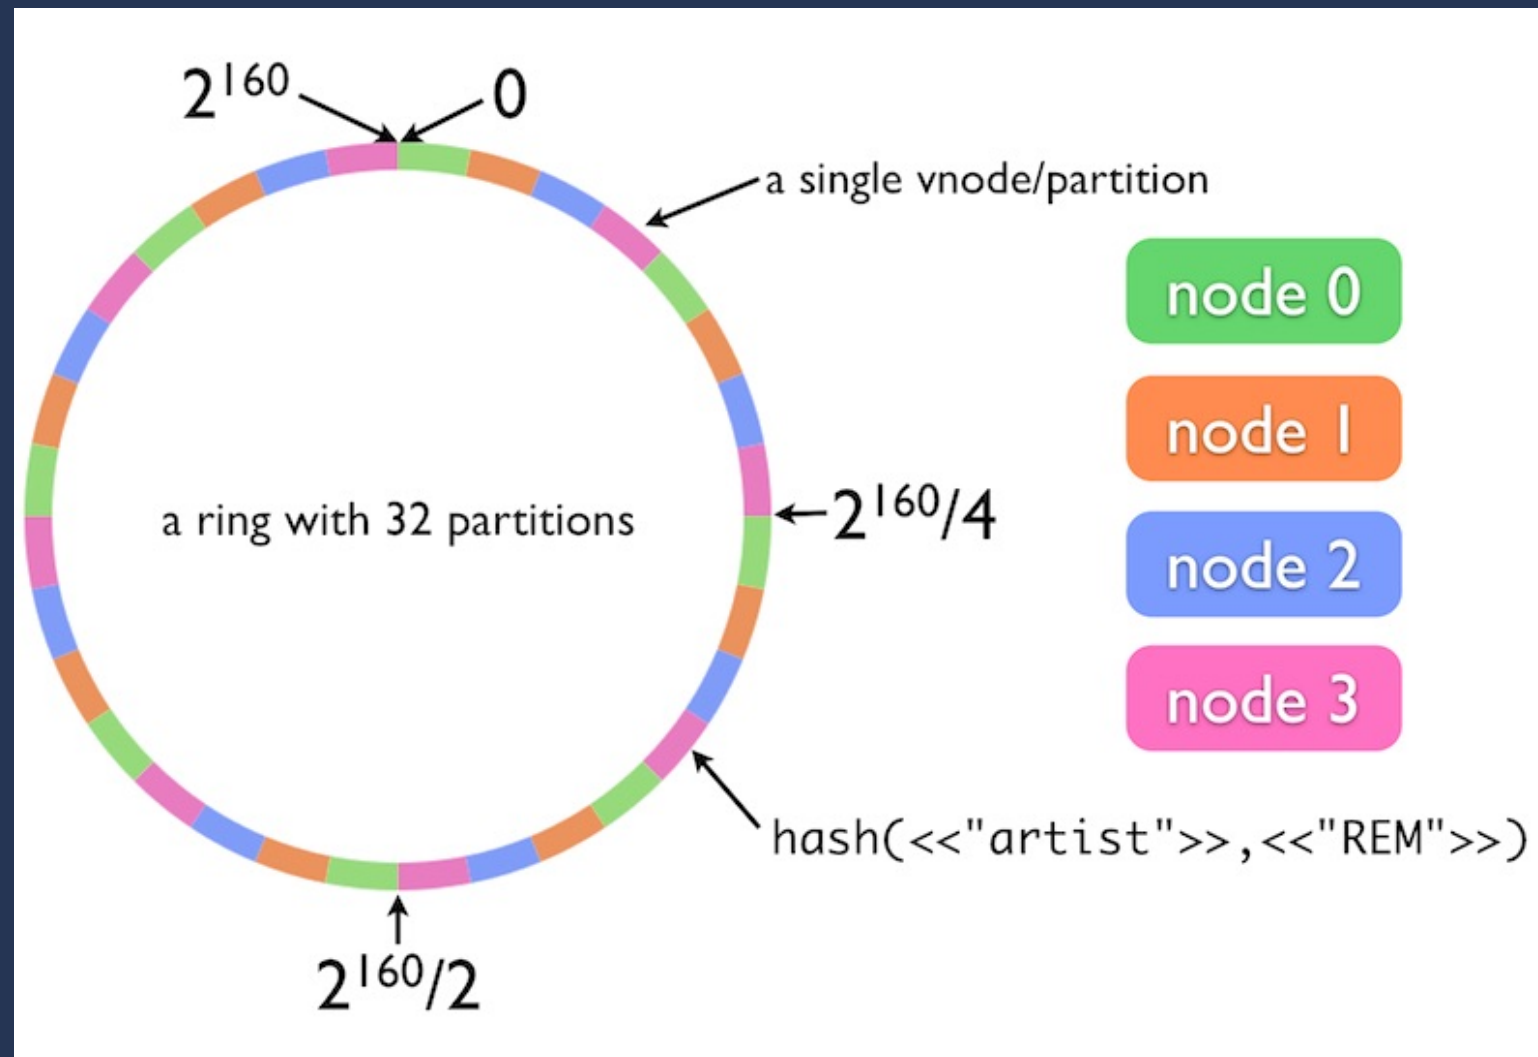
Requirement to keep the systems fault tolerant

- Redundancy: two or more resources for each unit of processing
- Supervising the failure of the units by an independent supervisor
- Rollback capability: undo the incomplete operations and retry

Supervisor



Consistent hashing of Basho Technologies' Riak database ⁶



⁶ (Note: Rikitake was a Basho Technologies employee during February to September 2013.)

<http://docs.basho.com/riak/latest/theory/concepts/Clusters/>

Fault tolerance of Riak


- Multiple copies for each data bucket
- Data evenly distributed to each cluster member node, more resilient to failures
- Even if a node fails, the other nodes respond with the valid data
- Recovery replication will happen after the node recovery
- *All automated*

The inconvenient truth of distributed systems

**Consistency:
hard or
impossible to
maintain**

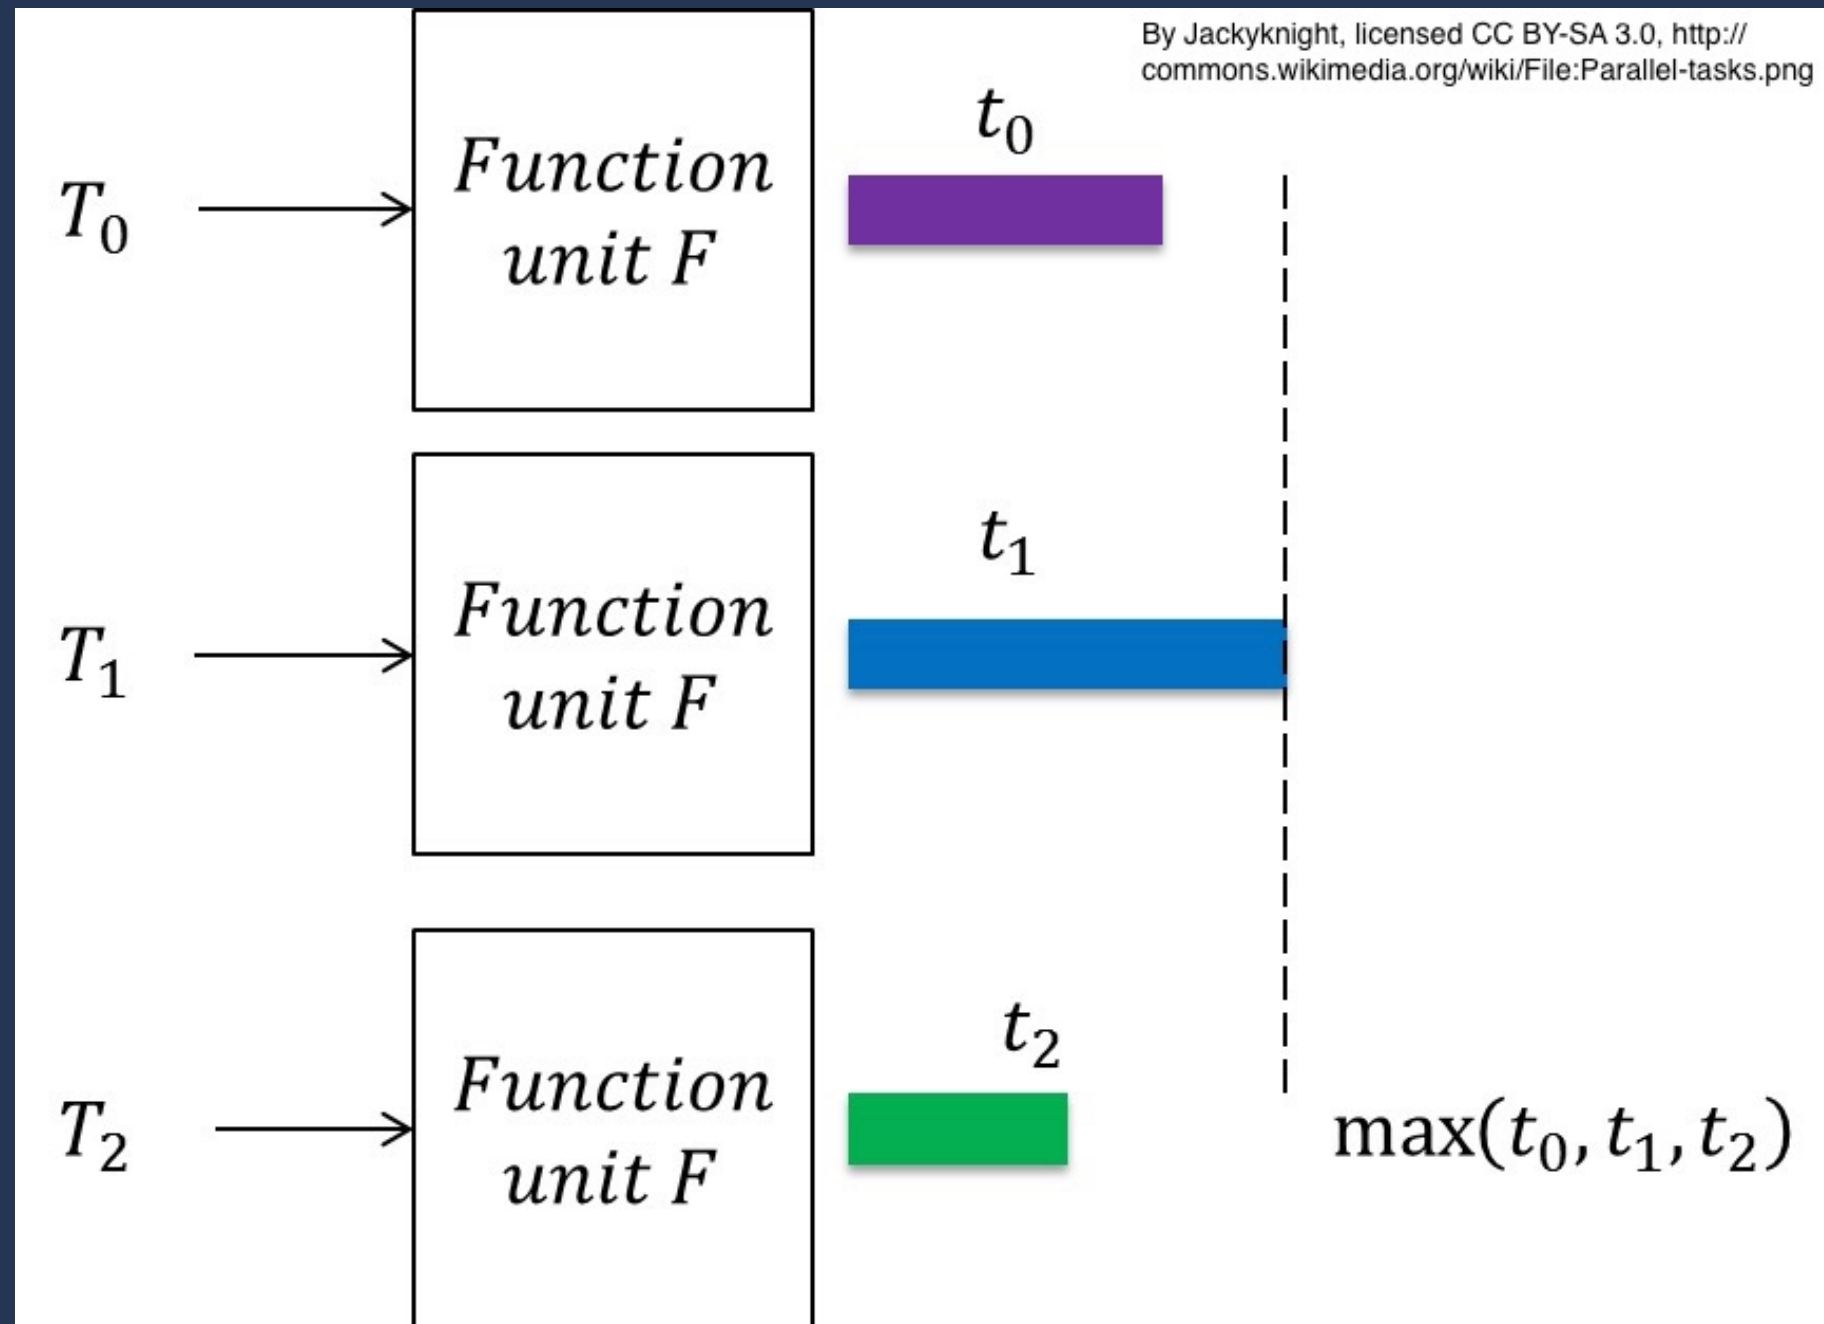
Net split

Recovery from net split is complex

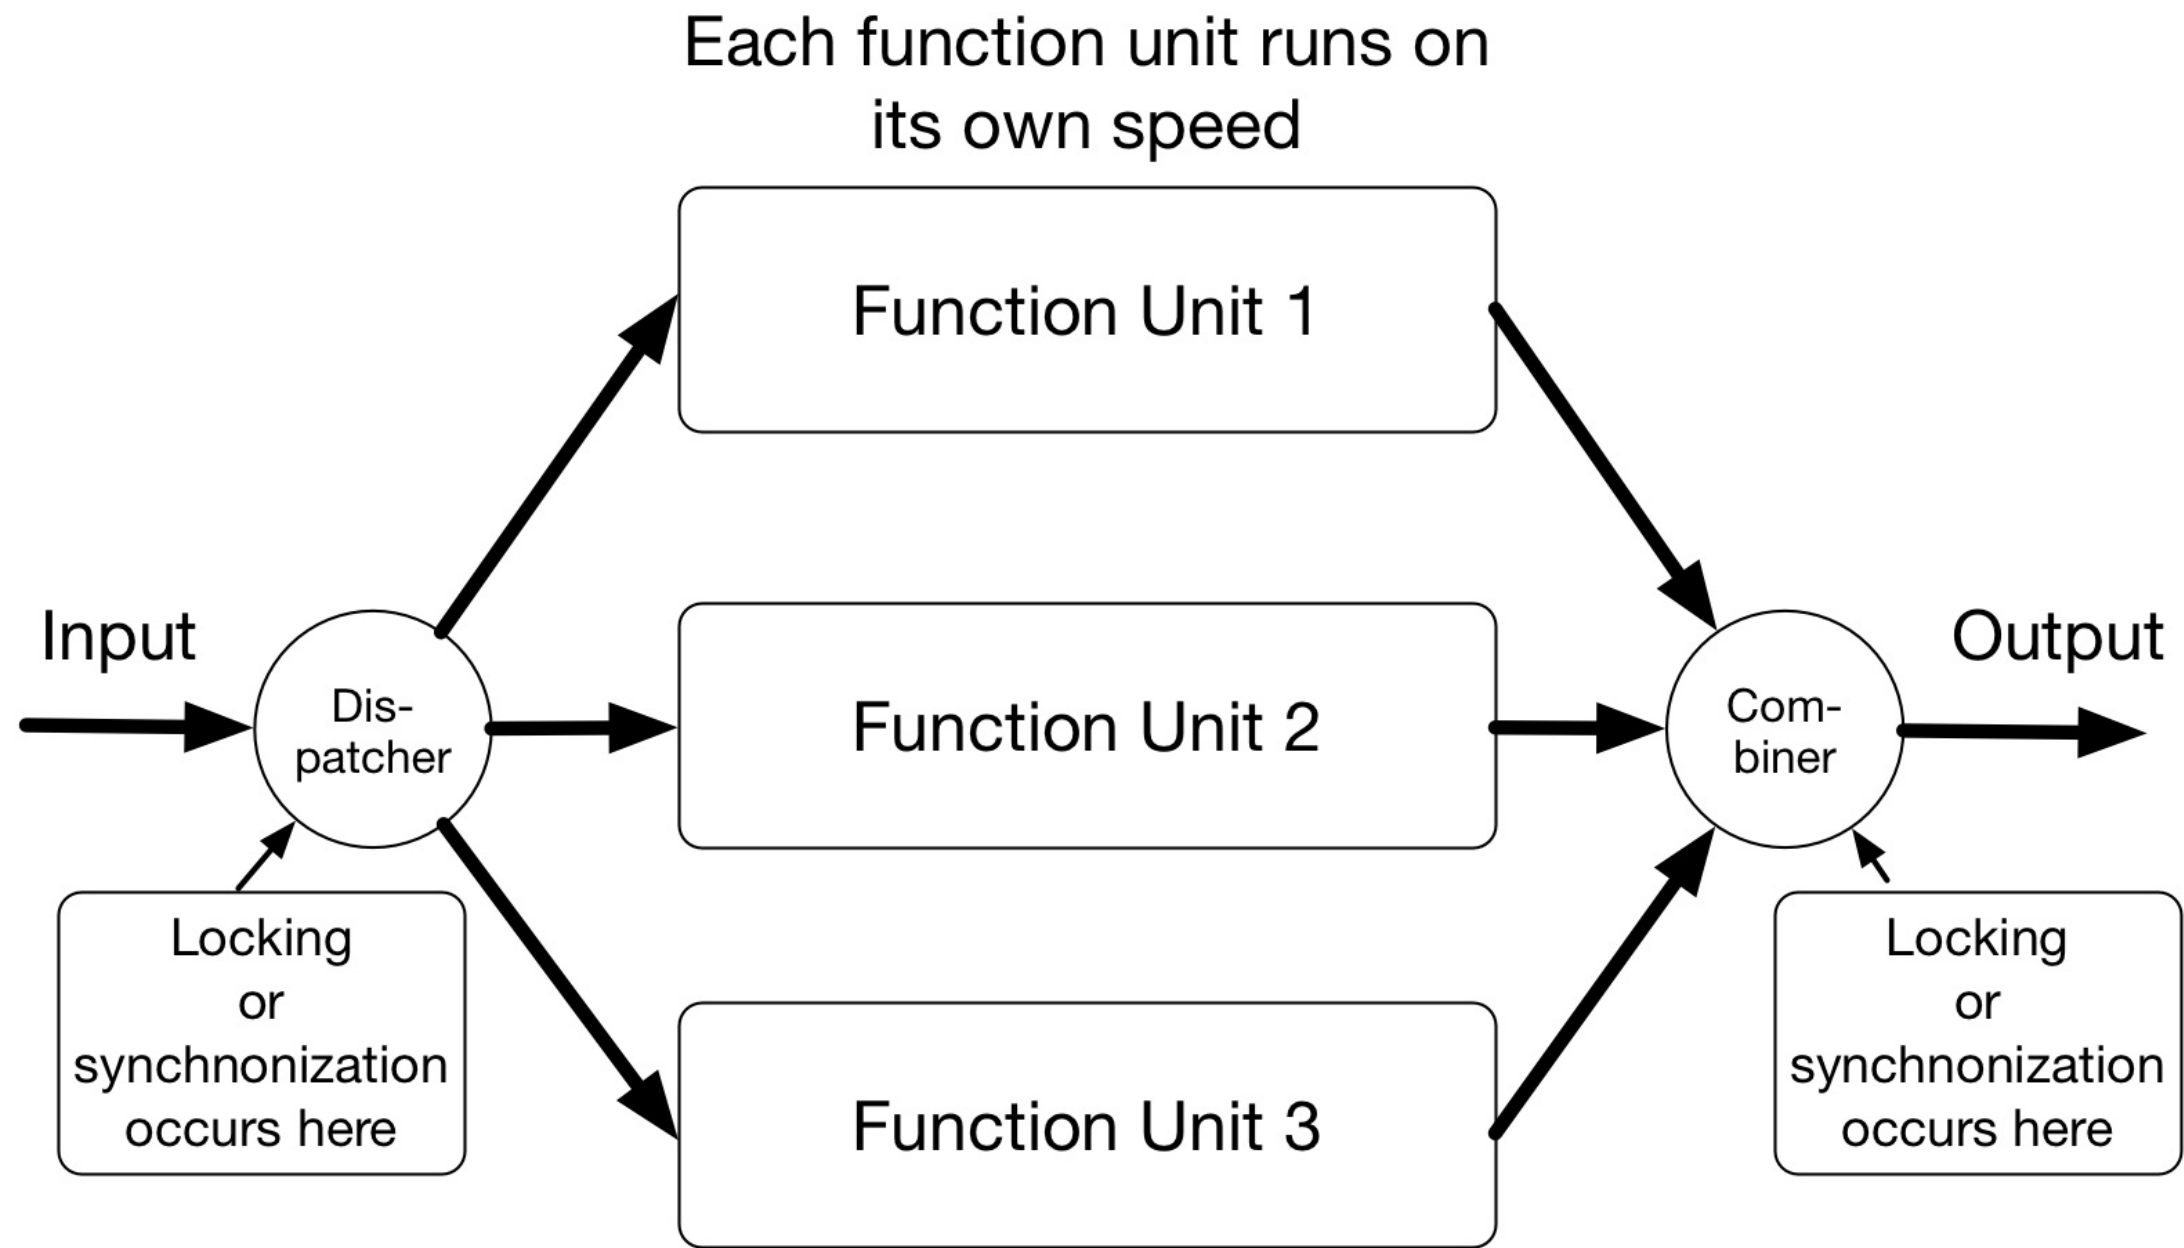


**Concurrency:
every system is
running on its
own;
synchronization
needed**

Synchronization



Locking



How precise the locking or synchronization timing should be?

- *It depends on the application*
- Bank transaction: strict
- Shopping cart: not necessarily strict
- Domain Name System: loose

**Current trend:
less locking,
more
inconsistency
allowance**

Questions

- How much is the cost of synchronization?
- Why do we need concurrent systems?
Stability? Performance?
- What have we traded in for obtaining the *convenient cloud computing systems*? Can we take them back?